

Familiarise yourself with the HMC sampling algorithm

Code to implement HMC is supplied in the module [hmc_Lab.py](#) and there is a simple demonstration of its usage in the notebook [demo_hmc.ipynb](#).

Included in [hmc_Lab.py](#) are the functions:

- [sample](#): The main function, to return samples from some arbitrary distribution over a state space (e.g. a posterior over combined model parameters θ)
- [one_step](#): Implements a single cycle of HMC (you needn't call this directly)
- [gradient_check](#): A useful helper to check your gradient calculations (can also be called automatically from within [sample](#)). Highly recommended!

Before undertaking any subsequent tasks, you should examine and run the example notebook [demo_hmc.ipynb](#), which offers a simple guide on to how to set up and use the sample algorithm.

The sampler is called as:

```
sample(x, energy_func, energy_grad, R, L, eps, burn=0, checkgrad=False, args=())
```

Arguments are:

- [x](#): The initial state (parameter values) of θ . In this exercise, any sensible random initialisation should be fine
- [energy_func](#): The energy (negative log-probability) of the distribution to be sampled from
- [energy_grad](#): Gradient of the above energy function with respect to θ
- [R](#): Number of samples desired (try a few hundred for testing, then perhaps up to 10,000 for quoting results)
- [L](#): Number of steps to run the Hamiltonian dynamics within each cycle of the sampler ($L=100$ is probably safe, though for the Gaussian example below, 25 should suffice)
- [eps](#): Step-size for the dynamics (see below)
- [burn](#): Number of initial samples to discard as 'burn-in' (e.g. you might set this to 10% of [R](#))
- [checkgrad](#): A flag specifying whether to check the consistency of the error/gradient functions – it is recommended to set this to [True](#)!
- [args](#): A list containing any additional (fixed) arguments that need to be passed to both [energy_func](#) and [energy_grad](#) (later, you will need to pass the data inputs and targets)

Key to the accuracy of the sampler are the 'simulation' parameters [L](#) and [eps](#). The former should be set large enough such that the dynamics have time to travel the full range of the distribution for a given step-size [eps](#), which itself should be sufficiently small to ensure the simulation does not 'diverge'. The best settings are rarely obvious without some experimentation!

Here, the general guidance is:

- Set **L** to 100 (as a minimum) for most regression exercises
- As an exception, setting **L** to 25 for the simple Gaussian example is sufficient
- Feel free to experiment and use a smaller **L** when testing
- Given choice of **L**, **eps** should be set to be as large as possible while keeping the acceptance rate of the sampler relatively high, e.g. no lower than 80%
- Note that having an acceptance rate of 100% is not likely to be a good thing. That may indicate that **eps** is too small and the state space is not being properly explored
- One informal approach may therefore be to increase **eps** until the acceptance rate falls below 80%, then 'dial it back down a little'.